

Chapter 1: Tic Tac Toe

Gabe Sekeres

February 17, 2025

Cornell University

Baseline Code

- There are advantages and disadvantages to any coding language

¹Matters a lot for LLMs!

- There are advantages and disadvantages to any coding language
- Python: extremely well-known, intuitive, better documented than anything else,¹ basically industry standard

¹Matters a lot for LLMs!

- There are advantages and disadvantages to any coding language
- Python: extremely well-known, intuitive, better documented than anything else,¹ basically industry standard
- Julia: none of these things, but extremely fast

¹Matters a lot for LLMs!

- There are advantages and disadvantages to any coding language
- Python: extremely well-known, intuitive, better documented than anything else,¹ basically industry standard
- Julia: none of these things, but extremely fast
- Zhang code: Python takes 37.33 seconds, Julia takes 3.29, and it reduces to 2.53 without printing

¹Matters a lot for LLMs!

- There are advantages and disadvantages to any coding language
- Python: extremely well-known, intuitive, better documented than anything else,¹ basically industry standard
- Julia: none of these things, but extremely fast
- Zhang code: Python takes 37.33 seconds, Julia takes 3.29, and it reduces to 2.53 without printing
- I use Julia. (I'll present the Python, it's more readable)

¹Matters a lot for LLMs!

Classes²

Classes are the objects of interest in this code. Each class has a number of inherent attributes, and essentially acts as a tuple of those attributes. We have four:

1. $\text{State} = \langle \text{data}, \text{winner}, \text{hash}, \text{end} \rangle$. This is a single board.
2. $\text{Player} = \langle \hat{V}(\mathcal{S}), \alpha, \varepsilon, \text{states}, \text{greedy}, \{X, O\} \rangle$. This contains all the parameters you'd expect, as well as two vectors, states and greedy.
3. $\text{Judger} = \langle P_1, P_2, \text{currentPlayer}, \text{currentState} \rangle$. This class runs the game. More on it later.
4. $\text{HumanPlayer} = \langle \{X, O\}, \text{state} \rangle$. This doesn't matter for us, I will be ignoring all of the human parts from here on.

²(in Julia, `struct`)

Class Functions

Each class has a set of functions inherent to it. They only use that class's elements. In Python, they are defined in the class, in Julia you define them elsewhere. Each class has a constructor function, where you call e.g. `State()` to build a state.

Remark. Best practice is [functional programming](#), for speed, modularity, and readability. The vast majority of what we'll do is defining functions of different types

State Functions - hash(State)

- Take in a certain state, and assign it a unique number (the hash value)

```
def hash(self):  
    if self.hash_val is None:  
        self.hash_val = 0  
    for i in np.nditer(self.data):  
        self.hash_val = self.hash_val * 3 + i + 1  
    return self.hash_val
```

State Functions - is_end(State) Part 1

- Take in a board, check if the game is over, if so add the winner and end value to the state, otherwise add false for both.

```
def is_end(self):
    if self.end is not None:
        return self.end
    results = []
    # check row
    for i in range(BOARD_ROWS):
        results.append(np.sum(self.data[i, :]))
    # check columns
    for i in range(BOARD_COLS):
        results.append(np.sum(self.data[:, i]))
    # check diagonals
    trace = 0
    reverse_trace = 0
    for i in range(BOARD_ROWS):
        trace += self.data[i, i]
        reverse_trace += self.data[i, BOARD_ROWS - 1 - i]
    results.append(trace)
    results.append(reverse_trace)
```

State Functions - is_end(State) Part 2

```
for result in results:
    if result == 3:
        self.winner = 1
        self.end = True
        return self.end
    if result == -3:
        self.winner = -1
        self.end = True
        return self.end

# whether it's a tie
sum_values = np.sum(np.abs(self.data))
if sum_values == BOARD_SIZE:
    self.winner = 0
    self.end = True
    return self.end

# game is still going on
self.end = False
return self.end
```

State Functions - next_state(State, i, j, symbol)

- Add a move to the board

```
def next_state(self, i, j, symbol):  
    new_state = State()  
    new_state.data = np.copy(self.data)  
    new_state.data[i, j] = symbol  
    return new_state
```

Global State Functions

- These get all of the states that are possible to attain from gameplay:

```
def get_all_states():
    current_symbol = 1
    current_state = State()
    all_states = dict()
    all_states[current_state.hash()] = (current_state, current_state.is_end())
    get_all_states_impl(current_state, current_symbol, all_states)
    return all_states
```

```
def get_all_states_impl(current_state, current_symbol, all_states):
    for i in range(BOARD_ROWS):
        for j in range(BOARD_COLS):
            if current_state.data[i][j] == 0:
                new_state = current_state.next_state(i, j, current_symbol)
                new_hash = new_state.hash()
                if new_hash not in all_states:
                    is_end = new_state.is_end()
                    all_states[new_hash] = (new_state, is_end)
                    if not is_end:
                        get_all_states_impl(new_state, -current_symbol,
                                            all_states)
```

Player Functions - reset(Player), set_state(Player, State)

- Return the player to the beginning of the game, resetting the attained states and their respective choices

```
def reset(self):  
    self.states = []  
    self.greedy = []
```

- Add a State to the list, with a greedy choice

```
def set_state(self, state):  
    self.states.append(state)  
    self.greedy.append(True)
```

Player Functions - set_symbol(Player, symbol)

- Add a symbol ($\{1, -1\} \equiv \{X, O\}$), and the initial estimations

```
def set_symbol(self, symbol):
    self.symbol = symbol
    for hash_val in all_states:
        state, is_end = all_states[hash_val]
        if is_end:
            if state.winner == self.symbol:
                self.estimated[hash_val] = 1.0
            elif state.winner == 0:
                # we need to distinguish between a tie and a lose
                self.estimated[hash_val] = 0.5
            else:
                self.estimated[hash_val] = 0
        else:
            self.estimated[hash_val] = 0.5
```


Player Functions - backup(Player)

- After each game, update the estimations using TD learning

```
def backup(self):
    states = [state.hash() for state in self.states]

    for i in reversed(range(len(states) - 1)):
        state = states[i]
        td_error = self.greedy[i] * (
            self.estimateds[states[i + 1]] - self.estimateds[state]
        )
        self.estimateds[state] += self.step_size * td_error
```

Math:

$$V(S_t) = V(S_t) + \alpha [V(S_{t+1}) - V(S_t)]$$

Player Functions - act(Player)

- Choose an action, based on $\text{rand}() \sim \text{Uniform}[0, 1)$ and ϵ

```
def act(self):
    state = self.states[-1]
    next_states = []
    next_positions = []
    for i in range(BOARD_ROWS):
        for j in range(BOARD_COLS):
            if state.data[i, j] == 0:
                next_positions.append([i, j])
                next_states.append(state.next_state(
                    i, j, self.symbol).hash())
    if np.random.rand() < self.epsilon:
        action = next_positions[np.random.randint(len(next_positions))]
        action.append(self.symbol)
        self.greedy[-1] = False
        return action
    values = []
    for hash_val, pos in zip(next_states, next_positions):
        values.append((self.estimated[hash_val], pos))
    # to select one of the actions of equal value at random due to Python's
    # sort is stable
    np.random.shuffle(values)
    values.sort(key=lambda x: x[0], reverse=True)
    action = values[0][1]
    action.append(self.symbol)
    return action
```

Player Functions - save_policy(Player) / load_policy(Player)

- Save the estimations and load them later. I did this very differently, where I had the train function return the converged expectations and kept them as variables. I don't understand why they did it like this.

```
def save_policy(self):  
    with open('policy_%s.bin' % ('first' if self.symbol == 1 else 'second'),  
              'wb') as f:  
        pickle.dump(self.estimations, f)  
  
def load_policy(self):  
    with open('policy_%s.bin' % ('first' if self.symbol == 1 else 'second'),  
              'rb') as f:  
        self.estimations = pickle.load(f)
```

The Judger Class

- Why does this exist?
- Short answer: functional programming
- Long answer: It's significantly easier to put this all in a different class, rather than having to alternate the players / define the parameters of the game manually.
- Tldr: It's slightly annoying but makes it easier to change the game. See Extensions below.

Judger Functions - reset(Judger) / alternate(Judger)

- Reset the two players

```
def reset(self):  
    self.p1.reset()  
    self.p2.reset()
```

- Switch who plays X and who plays O

```
def alternate(self):  
    while True:  
        yield self.p1  
        yield self.p2
```

Judger Functions - play(Judger)

- Run a single iteration of the game

```
def play(self):
    alternator = self.alternate()
    self.reset()
    current_state = State()
    self.p1.set_state(current_state)
    self.p2.set_state(current_state)
    while True:
        player = next(alternator)
        i, j, symbol = player.act()
        next_state_hash = current_state.next_state(i, j, symbol).hash()
        current_state, is_end = all_states[next_state_hash]
        self.p1.set_state(current_state)
        self.p2.set_state(current_state)
        if is_end:
            return current_state.winner
```

Global Functions - train(epochs, print_every_n=500)

- Play $N = \text{epochs}$ games, printing results every 500 iterations.
Have each player learn after each game.

```
def train(epochs, print_every_n=500):
    player1 = Player(epsilon=0.01)
    player2 = Player(epsilon=0.01)
    judger = Judger(player1, player2)
    player1_win = 0.0
    player2_win = 0.0
    for i in range(1, epochs + 1):
        winner = judger.play(print_state=False)
        if winner == 1:
            player1_win += 1
        if winner == -1:
            player2_win += 1
        if i % print_every_n == 0:
            print('Epoch %d, player 1 winrate: %.02f, player 2 winrate: %.02f'
                  % (i, player1_win / i, player2_win / i))
        player1.backup()
        player2.backup()
        judger.reset()
    player1.save_policy()
    player2.save_policy()
```

Global Functions - compete(turns)

- Play for turns games, where each player is always greedy

```
def compete(turns):  
    player1 = Player(epsilon=0)  
    player2 = Player(epsilon=0)  
    judger = Judger(player1, player2)  
    player1.load_policy()  
    player2.load_policy()  
    player1_win = 0.0  
    player2_win = 0.0  
    for _ in range(turns):  
        winner = judger.play()  
        if winner == 1:  
            player1_win += 1  
        if winner == -1:  
            player2_win += 1  
    judger.reset()  
    print('%d turns, player 1 win %.02f, player 2 win %.02f' % (turns,  
        player1_win / turns, player2_win / turns))
```


- Now that we have all the functions, this is all we need:

```
import numpy as np
import pickle

BOARD_ROWS = 3
BOARD_COLS = 3
BOARD_SIZE = BOARD_ROWS * BOARD_COLS
all_states = get_all_states()

train(int(1e5))
compete(int(1e3))
play()
```

Extensions

Questions


1. What happens if a tie is treated as a loss?
2. (Marco) What if we start at the analytic solution (henceforth, minimax estimations)?
3. What happens if we start with random estimations?
4. What happens if we train against a random player rather than a reinforcement learner?

- Since I used Julia, I was able to train for 3m epochs (with $\varepsilon = 0.1$) fairly easily. I initialized five different learners:
 1. Baseline, win worth 1, tie worth 0.5, loss worth 0
 2. Minimax, where we start at the minimax estimations
 3. Random, starting at random estimations
 4. No Ties, Baseline but ties worth 0
 5. Against Random, Baseline but trained against a random mover


Competition Results

- I played the models against each other for 10,000 rounds
- Baseline, Minimax, No Ties, and Against Random all were able to attain 100% tie rate (except when No Ties was playing O , of course)
- Random, where we drew each initial estimation from $\text{Uniform}[0, 1)$, lost every game against the other models

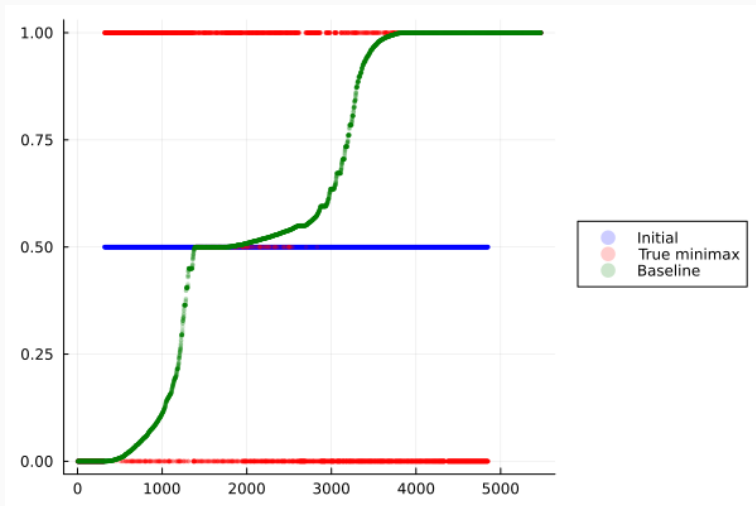
Estimation Differences (\times Player)




Estimation Differences (\times Player)




Estimation Differences (\times Player)




Estimation Differences (\times Player)




Estimation Differences (\times Player)




Estimation Differences (\times Player)




Estimation Differences (\times Player)




Estimation Differences (*O* Player)




Estimation Differences (*O* Player)




Estimation Differences (*O* Player)




Estimation Differences (*O* Player)



Estimation Differences (O Player)



Estimation Differences (*O* Player)



Estimation Differences (*O* Player)

