

# Reinforcement learning: An introduction

## Chapter 7

RL Reading Group

## Recap: MC & TD

- ▶ MC: Must wait until the end of the episode to determine the increment to  $V(S_t)$

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$

- ▶ TD: Form a target and make a useful update using the observed at time  $t + 1$

$$V(S_t) \leftarrow V(S_t) + \alpha \left[ \underbrace{R_{t+1} + \gamma V(S_{t+1})}_{\text{TD target } G_{t:t+1}} - V(S_t) \right]$$

- ▶ NEITHER ARE ALWAYS THE BEST!

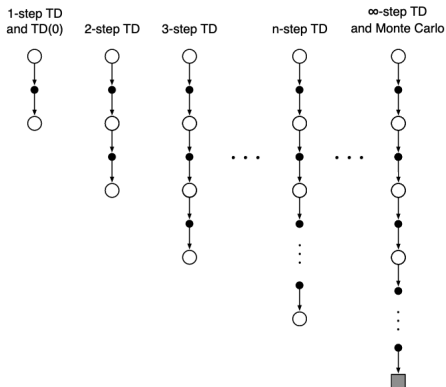
## Chapter 7: N-step TD

- ▶ target for a two-step update is the two-step return

$$G_{t:t+2} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

- ▶ target for a n-step update is the n-step return

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$$



## Chapter 7: N-step TD

- ▶ target for a two-step update is the two-step return

$$G_{t:t+2} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

- ▶ target for a n-step update is the n-step return

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$$

- ▶ The update rule is:

$$V_{t+n}(S_t) \leftarrow V_{t+n-1}(S_t) + \alpha [G_{t:t+n} - V_{t+n-1}(S_t)]$$

## Chapter 7: N-step TD

### **$n$ -step TD for estimating $V \approx v_\pi$**

Input: a policy  $\pi$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , a positive integer  $n$

Initialize  $V(s)$  arbitrarily, for all  $s \in \mathcal{S}$

All store and access operations (for  $S_t$  and  $R_t$ ) can take their index mod  $n + 1$

Loop for each episode:

    Initialize and store  $S_0 \neq \text{terminal}$

$T \leftarrow \infty$

    Loop for  $t = 0, 1, 2, \dots$ :

        If  $t < T$ , then:

            Take an action according to  $\pi(\cdot|S_t)$

            Observe and store the next reward as  $R_{t+1}$  and the next state as  $S_{t+1}$

            If  $S_{t+1}$  is terminal, then  $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$  ( $\tau$  is the time whose state's estimate is being updated)

        If  $\tau \geq 0$ :

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

            If  $\tau + n < T$ , then:  $G \leftarrow G + \gamma^n V(S_{\tau+n})$  ( $G_{\tau:\tau+n}$ )


$V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$

    Until  $\tau = T - 1$

# Chapter 7: Random Walk Example Revisited

## Example 6.2 Random Walk

In this example we empirically compare the prediction abilities of TD(0) and constant- $\alpha$  MC when applied to the following Markov reward process:



## Chapter 7: N-step Sarsa

- ▶ The target is:

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n})$$

- ▶ The update rule is:

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)]$$


## Chapter 7: N-step Sarsa

- ▶ The target is:

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n})$$

- ▶ The update rule is:

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)]$$



**Figure 7.4:** Gridworld example of the speedup of policy learning due to the use of  $n$ -step methods. The first panel shows the path taken by an agent in a single episode, ending at a location of high reward, marked by the G. In this example the values were all initially 0, and all rewards were zero except for a positive reward at G. The arrows in the other two panels show which action values were strengthened as a result of this path by one-step and  $n$ -step Sarsa methods. The one-step method strengthens only the last action of the sequence of actions that led to the high reward, whereas the  $n$ -step method strengthens the last  $n$  actions of the sequence, so that much more is learned from the one episode.



# Chapter 7: N-step Sarsa

## Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$   
Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$   
Loop for each episode:  
  Initialize  $S$   
  Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)  
  Loop for each step of episode:  
    Take action  $A$ , observe  $R, S'$   
    Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)  
     $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$   
     $S \leftarrow S'; A \leftarrow A';$   
  until  $S$  is terminal

## $n$ -step Sarsa for estimating $Q \approx q_*$ or $q_\pi$

Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}$   
Initialize  $\pi$  to be  $\varepsilon$ -greedy with respect to  $Q$ , or to a fixed given policy  
Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$ , a positive integer  $n$   
All store and access operations (for  $S_t$ ,  $A_t$ , and  $R_t$ ) can take their index mod  $n + 1$   
Loop for each episode:  
  Initialize and store  $S_0 \neq \text{terminal}$   
  Select and store an action  $A_0 \sim \pi(\cdot | S_0)$   
   $T \leftarrow \infty$   
  Loop for  $t = 0, 1, 2, \dots$ :  
    If  $t < T$ , then:  
      Take action  $A_t$   
      Observe and store the next reward as  $R_{t+1}$  and the next state as  $S_{t+1}$   
      If  $S_{t+1}$  is terminal, then:  
         $T \leftarrow t + 1$   
      else:  
        Select and store an action  $A_{t+1} \sim \pi(\cdot | S_{t+1})$   
     $\tau \leftarrow t - n + 1$  ( $\tau$  is the time whose estimate is being updated)  
    If  $\tau \geq 0$ :  
       $G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$   
      If  $\tau + n < T$ , then  $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$  ( $G_{\tau, \tau+n}$ )  
       $Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha [G - Q(S_\tau, A_\tau)]$   
      If  $\pi$  is being learned, then ensure that  $\pi(\cdot | S_\tau)$  is  $\varepsilon$ -greedy wrt  $Q$   
  Until  $\tau = T - 1$

## Chapter 7: N-step off-policy Sarsa

$$V_{t+n}(S_t) \leftarrow V_{t+n-1}(S_t) + \underbrace{\alpha}_{\text{importance sampling ratio}} \rho_{t:t+n-1} [G_{t:t+n} - V_{t+n-1}(S_t)]$$

### Off-policy $n$ -step Sarsa for estimating $Q \approx q_*$ or $q_{\pi}$

Input: an arbitrary behavior policy  $b$  such that  $b(a|s) > 0$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize  $\pi$  to be greedy with respect to  $Q$ , or as a fixed given policy

Algorithm parameters: step size  $\alpha \in (0, 1]$ , a positive integer  $n$

All store and access operations (for  $S_t$ ,  $A_t$ , and  $R_t$ ) can take their index mod  $n + 1$

Loop for each episode:

Initialize and store  $S_0 \neq \text{terminal}$

Select and store an action  $A_0 \sim b(\cdot|S_0)$

$$T \leftarrow \infty$$

Loop for  $t = 0, 1, 2, \dots$ :

- | If  $t < T$ , then:

Take action  $A_i$

Observe and st

If  $S_{t+1}$  is terminal, then:

$$T \leftarrow t + 1$$

```

|
|         else:

```

Select and store an action  $A_{t+1} \sim b(\cdot|S_{t+1})$

$$\tau \leftarrow t - n + 1 \quad (\tau \text{ is the time whose estimate is being updated})$$

If  $\tau \geq 0$ :

$$\rho \leftarrow \prod_{i=\tau+1}^{\min(\tau+n, T-1)} \frac{\pi(A_i|S_i)}{\pi(A_i|S_{i-1})} \quad (\rho_{\tau+1:\tau+n})$$
$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

If  $\tau + n \leq T$  then:  $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$

$$Q(S_{\pi}, A_{\pi}) \leftarrow Q(S_{\pi}, A_{\pi}) + \alpha \rho [G - Q(S_{\pi}, A_{\pi})]$$


If  $\pi$  is being learned, then ensure that  $\pi(\cdot|S_-)$  is greedy wrt  $Q$

Until  $\tau = T - 1$

## Chapter 7: Tree Backup

- The update rule is:

$$V_{t+n}(S_t) \leftarrow V_{t+n-1}(S_t) + \alpha [G_{t:t+n} - V_{t+n-1}(S_t)]$$



## Chapter 7: Tree Backup

- ▶ 1-step Expected Sarsa = 1-step Tree Backup

$$G_{t:t+1} = R_{t+1} + \sum_a \pi(a|S_{t+1})Q_t(S_{t+1}, a)$$

- ▶ 2-step Tree Backup

$$\begin{aligned} G_{t:t+2} &= R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1})Q_{t+1}(S_{t+1}, a) \\ &\quad + \gamma \pi(A_{t+1}|S_{t+1}) \left( R_{t+2} + \gamma \sum_a \pi(a|S_{t+2})Q_{t+1}(S_{t+2}, a) \right) \\ &= R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1})Q_{t+1}(S_{t+1}, a) \\ &\quad + \gamma \pi(A_{t+1}|S_{t+1})G_{t+1:t+2}, \end{aligned} \tag{1}$$

- ▶ n-step Tree Backup update (same as n-step Sarsa)

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)]$$

# Chapter 7: Tree Backup

## $n$ -step Tree Backup for estimating $Q \approx q_*$ or $q_\pi$

Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize  $\pi$  to be greedy with respect to  $Q$ , or as a fixed given policy

Algorithm parameters: step size  $\alpha \in (0, 1]$ , a positive integer  $n$

All store and access operations can take their index mod  $n + 1$

Loop for each episode:

    Initialize and store  $S_0 \neq \text{terminal}$

    Choose an action  $A_0$  arbitrarily as a function of  $S_0$ ; Store  $A_0$

$T \leftarrow \infty$

    Loop for  $t = 0, 1, 2, \dots$ :

        If  $t < T$ :

            Take action  $A_t$ ; observe and store the next reward and state as  $R_{t+1}, S_{t+1}$

            If  $S_{t+1}$  is terminal:

$T \leftarrow t + 1$

            else:

                Choose an action  $A_{t+1}$  arbitrarily as a function of  $S_{t+1}$ ; Store  $A_{t+1}$

$\tau \leftarrow t + 1 - n$  ( $\tau$  is the time whose estimate is being updated)

        If  $\tau \geq 0$ :

            If  $t + 1 \geq T$ :

$G \leftarrow R_T$

            else

$G \leftarrow R_{t+1} + \gamma \sum_a \pi(a|S_{t+1})Q(S_{t+1}, a)$

            Loop for  $k = \min(t, T - 1)$  down through  $\tau + 1$ :

$G \leftarrow R_k + \gamma \sum_{a \neq A_k} \pi(a|S_k)Q(S_k, a) + \gamma \pi(A_k|S_k)G$

$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha [G - Q(S_\tau, A_\tau)]$

            If  $\pi$  is being learned, then ensure that  $\pi(\cdot|S_\tau)$  is greedy wrt  $Q$

    Until  $\tau = T - 1$

# Chapter 7: Summary

